

Assignment - 4 Vajjah Srinivasa Taaran

In [1]:

```
#Import the Libraries.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
#Importing the dataset.
df = pd.read_csv("Employee-Attrition.csv")
```

In [3]:

```
df.head()
```

Out[3]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education- Field
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

In [4]:

```
df.shape
```

Out[4]:

```
(1470, 35)
```

In [5]:

```
df.Age.value_counts()
```

Out[5]:

35	78
34	77
36	69
31	69
29	68
32	61
30	60
33	58
38	58
40	57
37	50
27	48
28	48
42	46
39	42
45	41
41	40
26	39
44	33
46	33
43	32
50	30
25	26
24	26
49	24
47	24
55	22
51	19
53	19
48	19
54	18
52	18
22	16
56	14
23	14
58	14
21	13
20	11
59	10
19	9
18	8
60	5
57	4

Name: Age, dtype: int64

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object
22	OverTime	1470 non-null	object
23	PercentSalaryHike	1470 non-null	int64
24	PerformanceRating	1470 non-null	int64
25	RelationshipSatisfaction	1470 non-null	int64
26	StandardHours	1470 non-null	int64
27	StockOptionLevel	1470 non-null	int64
28	TotalWorkingYears	1470 non-null	int64
29	TrainingTimesLastYear	1470 non-null	int64
30	WorkLifeBalance	1470 non-null	int64
31	YearsAtCompany	1470 non-null	int64
32	YearsInCurrentRole	1470 non-null	int64
33	YearsSinceLastPromotion	1470 non-null	int64
34	YearsWithCurrManager	1470 non-null	int64

```
dtypes: int64(26), object(9)
```

```
memory usage: 402.1+ KB
```

In [7]:

```
df.describe()
```

Out[7]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNur
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.00
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.86
std	9.135373	403.509100	8.106864	1.024165	0.0	602.02
min	18.000000	102.000000	1.000000	1.000000	1.0	1.00
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.25
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.50
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.75
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.00

8 rows × 26 columns

In [8]:

```
#Checking for Null Values.  
df.isnull().any()
```

Out[8]:

Age	False
Attrition	False
BusinessTravel	False
DailyRate	False
Department	False
DistanceFromHome	False
Education	False
EducationField	False
EmployeeCount	False
EmployeeNumber	False
EnvironmentSatisfaction	False
Gender	False
HourlyRate	False
JobInvolvement	False
JobLevel	False
JobRole	False
JobSatisfaction	False
MaritalStatus	False
MonthlyIncome	False
MonthlyRate	False
NumCompaniesWorked	False
Over18	False
OverTime	False
PercentSalaryHike	False
PerformanceRating	False
RelationshipSatisfaction	False
StandardHours	False
StockOptionLevel	False
TotalWorkingYears	False
TrainingTimesLastYear	False
WorkLifeBalance	False
YearsAtCompany	False
YearsInCurrentRole	False
YearsSinceLastPromotion	False
YearsWithCurrManager	False

dtype: bool

In [9]:

```
df.isnull().sum()
```

Out[9]:

Age	0
Attrition	0
BusinessTravel	0
DailyRate	0
Department	0
DistanceFromHome	0
Education	0
EducationField	0
EmployeeCount	0
EmployeeNumber	0
EnvironmentSatisfaction	0
Gender	0
HourlyRate	0
JobInvolvement	0
JobLevel	0
JobRole	0
JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0
PercentSalaryHike	0
PerformanceRating	0
RelationshipSatisfaction	0
StandardHours	0
StockOptionLevel	0
TotalWorkingYears	0
TrainingTimesLastYear	0
WorkLifeBalance	0
YearsAtCompany	0
YearsInCurrentRole	0
YearsSinceLastPromotion	0
YearsWithCurrManager	0

dtype: int64

In [10]:

```
#Data Visualization.
```

```
sns.distplot(df["YearsWithCurrManager"])
```

```
/var/folders/0g/xqmh0yz92jx_s8ljsv3x08wr0000gn/T/ipykernel_82696/944934001.py:2: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(df["YearsWithCurrManager"])
```

Out[10]:

```
<Axes: xlabel='YearsWithCurrManager', ylabel='Density'>
```



In [11]:

df.corr()

/var/folders/0g/xqmh0yz92jx_s8ljsv3x08wr0000gn/T/ipykernel_82696/1134722465.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

df.corr()

Out[11]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber
Age	1.000000	0.010661	-0.001686	0.208034	NaN	NaN
DailyRate	0.010661	1.000000	-0.004985	-0.016806	NaN	NaN
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	NaN	NaN
Education	0.208034	-0.016806	0.021042	1.000000	NaN	NaN
EmployeeCount	NaN	NaN	NaN	NaN	NaN	NaN
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	NaN	NaN
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	NaN	NaN
HourlyRate	0.024287	0.023381	0.031131	0.016775	NaN	NaN
JobInvolvement	0.029820	0.046135	0.008783	0.042438	NaN	NaN
JobLevel	0.509604	0.002966	0.005303	0.101589	NaN	NaN
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	NaN	NaN
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	NaN	NaN
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	NaN	NaN
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	NaN	NaN
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	NaN	NaN
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	NaN	NaN
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	NaN	NaN
StandardHours	NaN	NaN	NaN	NaN	NaN	NaN
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	NaN	NaN
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	NaN	NaN
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	NaN	NaN
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	NaN	NaN
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	NaN	NaN
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	NaN	NaN
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	NaN	NaN
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	NaN	NaN

26 rows × 26 columns

In [12]:

```
df.head()
```

Out[12]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

In [13]:


```
plt.subplots(figsize = (25,25))
sns.heatmap(df.corr(),annot=True)
```

/var/folders/0g/xqmh0yz92jx_s8ljsv3x08wr0000gn/T/ipykernel_82696/3917469003.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True)
```

Out[13]:

<Axes: >




In [14]:

```
sns.boxplot(df.YearsWithCurrManager)
```

Out[14]:

<Axes: >



In [15]:

```
from scipy import stats
z_scores = np.abs(stats.zscore(df['YearsWithCurrManager']))
max_threshold=3
outliers = df['YearsWithCurrManager'][z_scores > max_threshold]

# Print and visualize the outliers
print("Outliers detected using Z-Score:")
print(outliers)
```

Outliers detected using Z-Score:

28	17
123	15
153	15
187	15
231	15
386	17
561	16
616	17
635	15
686	17
875	17
926	17
1078	17
1348	16

Name: YearsWithCurrManager, dtype: int64

In [16]:

```
q1 = df.YearsWithCurrManager.quantile(0.25)
q3 = df.YearsWithCurrManager.quantile(0.75)
print(q1)
print(q3)
upperlimit = q3+1.5*(q3-q1)
upperlimit
lowerlimit = q1-1.5*(q3-q1)
lowerlimit
df.median()
df["YearsWithCurrManager"] = np.where(df["YearsWithCurrManager"] > upperlimit, 14, df["YearsWithCurrManager"])
sns.boxplot(df.YearsWithCurrManager)
```

2.0


7.0

/var/folders/0g/xqmh0yz92jx_s8ljsv3x08wr0000gn/T/ipykernel_82696/1750964472.py:9: FutureWarning: The default value of numeric_only in DataFrame.median is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
df.median()
```

Out[16]:

<Axes: >



In [17]:

```

from scipy import stats
z_scores = np.abs(stats.zscore(df['YearsWithCurrManager']))
max_threshold=3
outliers = df['YearsWithCurrManager'][z_scores > max_threshold]

# Print and visualize the outliers
print("Outliers detected using Z-Score:")
print(outliers)

```

Outliers detected using Z-Score:
 Series([], Name: YearsWithCurrManager, dtype: int64)

In [18]:

```
df.head()
```

Out[18]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

In [19]:

```

x=df.drop('Attrition',axis=1)
x.head()

```

Out[19]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences
2	37	Travel_Rarely	1373	Research & Development	2	2	Other
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences
4	27	Travel_Rarely	591	Research & Development	2	1	Medical

5 rows × 34 columns

In [20]:

```
y=df.Attrition
y.head()
```

Out[20]:

```
0    Yes
1    No
2    Yes
3    No
4    No
Name: Attrition, dtype: object
```

In [21]:

```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
x.BusinessTravel =le.fit_transform(x.BusinessTravel )
x.head()
x.Department =le.fit_transform(x.Department )
x.head()
x.EducationField =le.fit_transform(x.EducationField )
x.head()
x.Gender=le.fit_transform(x.Gender)
x.head()
x.JobRole =le.fit_transform(x.JobRole )
x.head()
x.MaritalStatus =le.fit_transform(x.MaritalStatus )
x.head()
x.Over18 =le.fit_transform(x.Over18 )
x.head()
x.OverTime =le.fit_transform(x.OverTime )
x.head()
```

Out[21]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	E
0	41	2	1102	2	1	2	1	
1	49	1	279	1	8	1	1	
2	37	2	1373	1	2	2	4	
3	33	1	1392	1	3	4	1	
4	27	2	591	1	2	1	3	

5 rows × 34 columns

In [22]:

df.columns

Out[22]:

```
Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Departmen
t',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCou
nt',
      'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyR
ate',
      'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
      'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesW
orked',
      'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLeve
l',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalanc
e',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotio
n',
      'YearsWithCurrManager'],
      dtype='object')
```

In [23]:

```
#feature scaling
from sklearn.preprocessing import MinMaxScaler
ms=MinMaxScaler()
x_scaled=pd.DataFrame(ms.fit_transform(x),columns=x.columns)
```

In [24]:

x_scaled

Out[24]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educator
0	0.547619	1.0	0.715820	1.0	0.000000	0.25	
1	0.738095	0.5	0.126700	0.5	0.250000	0.00	
2	0.452381	1.0	0.909807	0.5	0.035714	0.25	
3	0.357143	0.5	0.923407	0.5	0.071429	0.75	
4	0.214286	1.0	0.350036	0.5	0.035714	0.00	
...	
1465	0.428571	0.5	0.559771	0.5	0.785714	0.25	
1466	0.500000	1.0	0.365784	0.5	0.178571	0.00	
1467	0.214286	1.0	0.037938	0.5	0.107143	0.50	
1468	0.738095	0.5	0.659270	1.0	0.035714	0.50	
1469	0.380952	1.0	0.376521	0.5	0.250000	0.50	

1470 rows × 34 columns

In [25]:

```
#Splitting Data into Train and Test.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2,random_state
```

In [26]:

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[26]:

```
((1176, 34), (294, 34), (1176,), (294,))
```

In [27]:

```
x_train.head()
```

Out[27]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educator
1374	0.952381	1.0	0.360057	1.0	0.714286	0.50	
1092	0.642857	1.0	0.607015	0.5	0.964286	0.50	
768	0.523810	1.0	0.141732	1.0	0.892857	0.50	
569	0.428571	0.0	0.953472	1.0	0.250000	0.75	
911	0.166667	0.5	0.355762	1.0	0.821429	0.00	

5 rows × 34 columns

In [28]:

```
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
```



```
model.fit(x_train,y_train)
pred=model.predict(x_test)
pred
```

```
array([ 'No', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'No', 'No', 'No',
'No',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No',
'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'Yes', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Ye
s',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'N
o',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'Yes', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No' ],
dtype=object)
```

In [30]:

```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
y=le.fit_transform(y)
```

In [31]:

```
y_test
```

Out[31]:

```
442      No
1091     No
981      Yes
785      No
1332     Yes
...
1439     No
481      No
124      Yes
198      No
1229     No
Name: Attrition, Length: 294, dtype: object
```

In [32]:

```
df
```

Out[32]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Ec
0	41	Yes	Travel_Rarely	1102	Sales	1	2	
1	49	No	Travel_Frequently	279	Research & Development	8	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	
4	27	No	Travel_Rarely	591	Research & Development	2	1	
...	
1465	36	No	Travel_Frequently	884	Research & Development	23	2	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	
1468	49	No	Travel_Frequently	1023	Sales	2	3	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	

1470 rows × 35 columns

Evaluation of classification model

In [33]:

```
#Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score
```

In [34]:

```
accuracy_score(y_test,pred)
```

Out[34]:

0.8843537414965986

In [35]:

```
confusion_matrix(y_test,pred)
```

Out[35]:

```
array([[242,  3],
       [ 31, 18]])
```

In [36]:

```
pd.crosstab(y_test,pred)
```

Out[36]:

col_0	No	Yes
Attrition		
No	242	3
Yes	31	18

Roc-AUC curve

In [37]:

```
probability=model.predict_proba(x_test)[:,-1]  
probability
```

Out[37]:

```
array([0.15843867, 0.20617997, 0.31691729, 0.09672152, 0.63876647,
       0.06205401, 0.61414184, 0.07466397, 0.00797252, 0.39157785,
       0.05281564, 0.33160211, 0.02022395, 0.6671328 , 0.19419683,
       0.0335299 , 0.10954936, 0.17130578, 0.043804 , 0.2241511 ,
       0.23531373, 0.01475346, 0.06562592, 0.05019163, 0.59115162,
       0.44667993, 0.07401303, 0.0449937 , 0.67637047, 0.05859033,
       0.01545736, 0.03386798, 0.07021403, 0.1707141 , 0.07767295,
       0.04154894, 0.08312937, 0.06997437, 0.03567429, 0.05269126,
       0.05742727, 0.02144976, 0.01779053, 0.01301572, 0.02825292,
       0.50162054, 0.41541766, 0.00299378, 0.74315718, 0.51799699,
       0.09708281, 0.48942319, 0.07941138, 0.25720931, 0.66861063,
       0.26482373, 0.01970983, 0.30281497, 0.02858501, 0.16213966,
       0.02040161, 0.2173984 , 0.13768821, 0.03568054, 0.37558052,
       0.03010741, 0.29718154, 0.15832399, 0.10264349, 0.08700774,
       0.0815183 , 0.30943969, 0.08708969, 0.07442596, 0.12300414,
       0.0618342 , 0.04633075, 0.07672219, 0.19834226, 0.03129952,
       0.00857215, 0.02394842, 0.13606932, 0.02587787, 0.03217004,
       0.0821409 , 0.00518749, 0.035308 , 0.03813342, 0.14270872,
       0.26418695, 0.16461435, 0.27401734, 0.24146954, 0.02119787,
       0.17774284, 0.34102562, 0.28338745, 0.06906981, 0.04948532,
       0.24465264, 0.74929682, 0.35691434, 0.01878265, 0.08772637,
       0.03239915, 0.05413857, 0.15215059, 0.07127406, 0.13828798,
       0.09342465, 0.04693869, 0.02494493, 0.15041914, 0.07133392,
       0.03025642, 0.05306455, 0.1165452 , 0.00872431, 0.01229042,
       0.17575238, 0.05005249, 0.09018395, 0.82857166, 0.03066995,
       0.0228189 , 0.00874605, 0.13496234, 0.16593413, 0.05060052,
       0.01520085, 0.29791945, 0.54919611, 0.33581407, 0.0469494 ,
       0.38773566, 0.61348127, 0.14171081, 0.07455884, 0.2409655 ,
       0.09528764, 0.06730943, 0.09797576, 0.20026612, 0.20053142,
       0.03046036, 0.14877431, 0.0036571 , 0.11146887, 0.15912883,
       0.06017571, 0.17964687, 0.06063618, 0.1199213 , 0.03284092,
       0.02688355, 0.06536903, 0.08335812, 0.01464284, 0.01536292,
       0.37701597, 0.01262506, 0.15004068, 0.80530948, 0.11655522,
       0.28461049, 0.17042029, 0.15392139, 0.02756879, 0.00599553,
       0.04142216, 0.09958411, 0.11567269, 0.10448555, 0.01830036,
       0.1444171 , 0.1048541 , 0.10079777, 0.05099176, 0.09183576,
       0.02893646, 0.09754427, 0.00516687, 0.75206394, 0.04227453,
       0.04018918, 0.37563319, 0.04457964, 0.72551665, 0.10583031,
       0.36656526, 0.38293703, 0.32923777, 0.05248015, 0.08216713,
       0.13748888, 0.04309097, 0.01429957, 0.2656631 , 0.06297408,
       0.16075744, 0.15388494, 0.67190498, 0.05834473, 0.28467369,
       0.04694404, 0.46237195, 0.00339026, 0.13927388, 0.02695884,
       0.12707414, 0.17395277, 0.0750947 , 0.10135673, 0.16496216,
       0.02583798, 0.01790826, 0.08850395, 0.02838351, 0.13795992,
       0.08655223, 0.22164621, 0.73379009, 0.17294814, 0.40907888,
       0.01503347, 0.11411826, 0.21412683, 0.32566668, 0.03366086,
       0.04472831, 0.32127248, 0.05442236, 0.0242917 , 0.16228044,
       0.32858438, 0.22879119, 0.00852736, 0.0798162 , 0.01140248,
       0.14102568, 0.29116266, 0.01282151, 0.17118076, 0.04051376,
       0.04165738, 0.42684273, 0.35009936, 0.0366853 , 0.11692325,
       0.37940034, 0.31562415, 0.79587005, 0.05488792, 0.21568794,
       0.06397987, 0.00569145, 0.66085682, 0.35796045, 0.37592133,
       0.3650533 , 0.03568965, 0.21192376, 0.05892118, 0.06428028,
       0.10143977, 0.00796354, 0.2678938 , 0.4288445 , 0.0652538 ,
       0.09309022, 0.01226927, 0.14314823, 0.04989664, 0.02304292,
       0.02508766, 0.06618985, 0.24272596, 0.26663754, 0.1979951 ,
       0.26504226, 0.01648205, 0.15826843, 0.08519882, 0.02669729,
       0.18757572, 0.00768502, 0.27928747, 0.0027473 , 0.02506718,
       0.22608608, 0.72428674, 0.07739605, 0.26575953])
```

In [38]:


```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
y_test=le.fit_transform(y_test)
```

In [39]:

```
# roc_curve
fpr,tpr,thresholds = roc_curve(y_test,probability)
```

In [40]:

```
plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



DecisionTreeClassifier

In [41]:

```
from sklearn.tree import DecisionTreeClassifier
dtt=DecisionTreeClassifier()
```

In [42]:

```
dtc.fit(x_train,y_train)
```

Out[42]:

```
▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

In [43]:

```
pred=dtc.predict(x_test)
```

In [44]:

pred

Out[44]:

```

array(['No', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'No', 'No', 'No',
      'No',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes',
      'No',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'Y
es',
      'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'Y
es',
      'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No',
      'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No',
      'No',
      'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
      'No',
      'No', 'Yes', 'Yes', 'Yes', 'No', 'No', 'No', 'No', 'Yes', 'No',
      'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No',
      'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'N
o',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'N
o',
      'No', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
      'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No',
      'No',
      'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No',
      'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'N
o',
      'No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
      'No',
      'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No', 'No',
      'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Ye
s',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'Yes',
      'No',
      'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'Y
es',
      'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
      'No',
      'No',
      'No', 'No', 'No' ], dtype=object)

```


In [45]:

```
y_test
```

Out[45]:

```
array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
      0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
      1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1,
      1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
1,
      0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
      0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
      1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
      0, 1, 0, 0, 0, 1, 0, 0])
```

In [46]:

df

Out[46]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Ec
0	41	Yes	Travel_Rarely	1102	Sales	1	2	
1	49	No	Travel_Frequently	279	Research & Development	8	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	
4	27	No	Travel_Rarely	591	Research & Development	2	1	
...
1465	36	No	Travel_Frequently	884	Research & Development	23	2	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	
1468	49	No	Travel_Frequently	1023	Sales	2	3	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	

1470 rows × 35 columns

Evaluation of classification model

In [47]:

```
#Accuracy score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score
```

In [48]:

```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
y=le.fit_transform(y)
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
pred=le.fit_transform(pred)
```

In [49]:

```
accuracy_score(y_test,pred)
```

Out[49]:

```
0.7721088435374149
```

In [50]:

```
confusion_matrix(y_test,pred)
```

Out[50]:

```
array([[211,  34],
       [ 33,  16]])
```

In [51]:

```
pd.crosstab(y_test,pred)
```

Out[51]:

	col_0	0	1
row_0			
0	211	34	
1	33	16	

In [52]:

```
print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.86	0.86	0.86	245
1	0.32	0.33	0.32	49
accuracy			0.77	294
macro avg	0.59	0.59	0.59	294
weighted avg	0.77	0.77	0.77	294

Roc-AUC curve

In [53]:

```
probability=dtc.predict_proba(x_test)[:,-1]
```

In [54]:

```
probability
```

Out[54]:


```
array([0., 0., 0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0.,
0.,
      0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.,
0.,
      0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0.,
1.,
      1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 0., 0.,
0.,
      0., 0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
0.,
      0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.,
0.,
      1., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.,
0.,
      0., 0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 0., 1., 1., 1., 0.,
0.,
      0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0.,
0.,
      0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0.,
0.,
      0., 0., 0., 0., 0., 1., 0., 1., 1., 0., 0., 0., 0., 1., 0., 0.,
1.,
      0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
0.,
      1., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
0.,
      0., 1., 0., 0., 0., 0., 1., 1., 0., 0., 1., 0., 0., 0., 0., 0.,
0.,
      0., 1., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
0.,
      0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 0., 0., 0.,
0.,
      0., 0., 0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 0., 1., 0., 0.,
0.,
      0., 0., 0., 0., 0.]
```

In [55]:

```
fpr, tpr, thresholds = roc_curve(y_test, probability)
```

In [56]:

```
plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



In [57]:

```
from sklearn import tree
plt.figure(figsize=(25,15))
tree.plot_tree(dtc,filled=True)
```

Out[57]:

```
[Text(0.33072029386712093, 0.9722222222222222, 'x[27] <= 0.038\ngini =
0.269\nsamples = 1176\nvalue = [988, 188]'),
Text(0.0817717206132879, 0.9166666666666666, 'x[16] <= 0.75\ngini =
0.5\nsamples = 78\nvalue = [39, 39]'),
Text(0.05110732538330494, 0.8611111111111112, 'x[4] <= 0.554\ngini =
0.426\nsamples = 39\nvalue = [27, 12]'),
Text(0.034071550255536626, 0.8055555555555556, 'x[15] <= 0.167\ngini
= 0.312\nsamples = 31\nvalue = [25, 6]'),
Text(0.020442930153321975, 0.75, 'x[21] <= 0.5\ngini = 0.49\nsamples
= 7\nvalue = [3, 4]'),
Text(0.013628620102214651, 0.6944444444444444, 'x[16] <= 0.25\ngini =
0.375\nsamples = 4\nvalue = [3, 1]'),
Text(0.0068143100511073255, 0.6388888888888888, 'gini = 0.0\nsamples
= 3\nvalue = [3, 0]'),
Text(0.020442930153321975, 0.6388888888888888, 'gini = 0.0\nsamples =
1\nvalue = [0, 1]'),
Text(0.027257240204429302, 0.6944444444444444, 'gini = 0.0\nsamples =
3\nvalue = [0, 31]')]
```

In [58]:

```
from sklearn.model_selection import GridSearchCV
parameter={
    'criterion':['gini','entropy'],
    'splitter':['best','random'],
    'max_depth':[1,2,3,4,5],
    'max_features':['auto','sqrt','log2']
}
```

In [59]:

```
grid_search=GridSearchCV(estimator=dtc,param_grid=parameter,cv=5,scoring="accuracy")
```

In [60]:

```
grid_search.fit(x_train,y_train)
```

```
/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/model_selection/_validation.py:425: FitFailedWarning:
100 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by setting
error_score='raise'.
```

Below are more details about the failures:


```
-----
100 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/model_selection/_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/base.py", line 1144, in wrapper
    estimator._validate_params()
  File "/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/base.py", line 637, in _validate_params
    validate_parameter_constraints(
  File "/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/utils/_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of DecisionTreeClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/taaran/anaconda3/lib/python3.11/site-packages/sklearn/model_selection/_search.py:976: UserWarning: One or more of the test scores are
non-finite: [          nan          nan 0.84013704 0.84013704 0.84013704 0.84013704
```

```
nan          nan 0.8384313  0.83843491 0.82311576 0.84013704
nan          nan 0.84437793 0.83928597 0.8452326  0.83758745
nan          nan 0.83759106 0.84183916 0.83161558 0.83928237
nan          nan 0.84608006 0.83673278 0.83163     0.84354129
nan          nan 0.84013704 0.84013704 0.84013704 0.84013704
nan          nan 0.83759466 0.84013704 0.83674721 0.84013704
nan          nan 0.84267941 0.84013704 0.83841688 0.83928597
nan          nan 0.83588532 0.84183916 0.84012261 0.83589975
nan          nan 0.82824739 0.83589254 0.82483231 0.83843491]
```

```
warnings.warn(
```

Out[60]:



In [61]:

```
grid_search.best_params_
```

Out[61]:

```
{'criterion': 'gini',
 'max_depth': 5,
 'max_features': 'sqrt',
 'splitter': 'best'}
```

In [62]:

```
dtc_cv=DecisionTreeClassifier(criterion= 'entropy',
                             max_depth=3,
                             max_features='sqrt',
                             splitter='best')
dtc_cv.fit(x_train,y_train)
```

Out[62]:

```
▼                               DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3, max_features
='sqrt')
```

In [63]:

```
pred=dtc_cv.predict(x_test)
```

In [64]:

```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
y=le.fit_transform(y)
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
pred=le.fit_transform(pred)
```

In [65]:

```
print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.84	1.00	0.91	245
1	1.00	0.04	0.08	49
accuracy			0.84	294
macro avg	0.92	0.52	0.50	294
weighted avg	0.87	0.84	0.77	294

RandomForestClassifier

In [66]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
```

In [67]:

```
forest_params = [{'max_depth': list(range(10, 15)), 'max_features': list(range(0,14))}]
```

In [68]:

```
rfc_cv= GridSearchCV(rfc,param_grid=forest_params,cv=10,scoring="accuracy")
```

In []:

```
rfc_cv.fit(x_train,y_train)
```

In []:

```
pred=rfc_cv.predict(x_test)
```

In []:

```
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
y=le.fit_transform(y)
#label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
pred=le.fit_transform(pred)
```

In []:

```
print(classification_report(y_test,pred))
```

In []:

```
rfc_cv.best_params_
```